


Docker

Jesús Morán y Cristian Augusto

Grupo de Investigación en Ingeniería del Software

<http://giis.uniovi.es>

Universidad de Oviedo



Problemas de “desarrollo”

No sé qué pudo ocurrir... en mi PC funcionaba

Instala la nueva versión de java


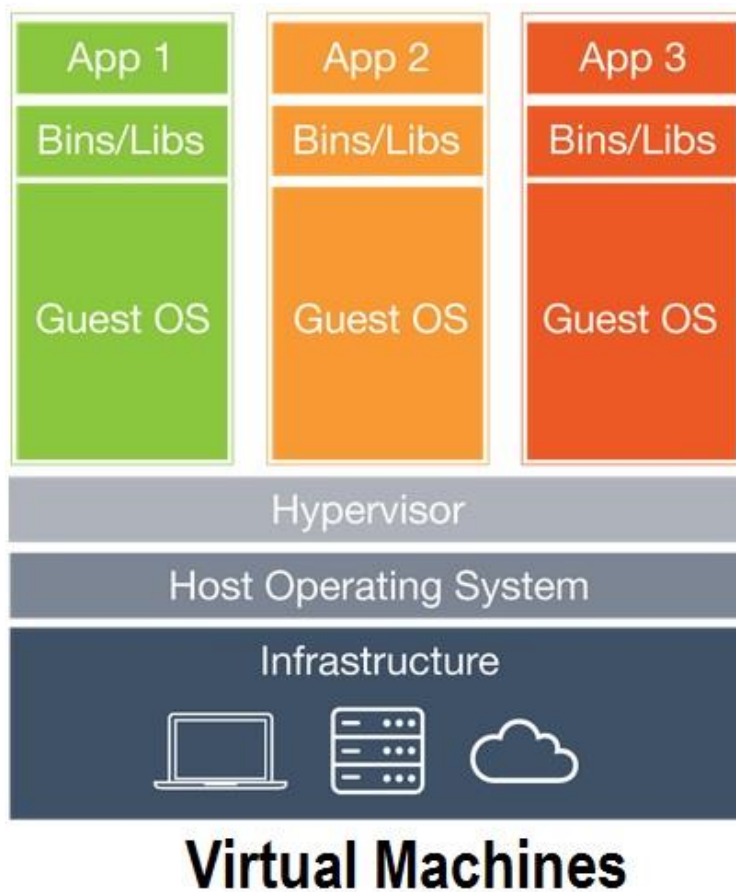
Falta una librería

Tienes que abrir el puerto 8000

Prueba a instalar tomcat

...

Máquina virtual vs Docker



Crédito de la imagen: docker.com

Terminología docker

- **Imagen:** *“An executable package that includes everything needed to run an application -the code, a runtime, libraries, environment variables, and configuration files.”* Docker.com
- **Contenedor:** *“A runtime instance of an image--what the image becomes in memory when executed (that is, an image with state, or a user process”* Docker.com

Docker Compose

■ Instalación:

☐ Descargar docker-compose

```
jesus@injtest:/disk0$ sudo curl -L "https://github.com/docker/compose/releases/download/1.29.2/docker-compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
[sudo] password for jesus:
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 664      100 664    0     0   5487      0 --:--:-- --:--:-- --:--:--   5487
100 12.1M    100 12.1M    0     0  15.9M      0 --:--:-- --:--:-- --:--:--  68.2M
```

☐ Dar permisos de ejecución

```
jesus@injtest:/disk0$ sudo chmod +x /usr/local/bin/docker-compose
```

☐ Instalar command line completion

```
jesus@injtest:/disk0$ sudo curl -L https://raw.githubusercontent.com/docker/compose/1.29.2/contrib/completion/bash/docker-compose -o /etc/bash_completion.d/docker-compose
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 13500    100 13500    0     0   117k      0 --:--:-- --:--:-- --:--:--   117k
```

☐ Actualizar la configuración de la terminal

```
jesus@injtest:/disk0$ source ~/.bashrc
```

Docker Compose

```
docker-compose --version
```

- Desde Docker Desktop 3.4.0 hay dos versiones
 - La versión docker-compose: implementada en python
 - La versión docker compose (dentro de docker-cli):
 - Se conoce como Compose V2
 - Implementada en Go
 - Todavía no tiene toda la funcionalidad de docker-compose
 - Más amigable con la nube:
 - Backends para guardar imágenes: ECS de AWS, ACI de Azure,...

```
docker compose --version
```

Docker Compose

```
version: '3.3'

services:
  db:
    image: mysql:5.7
    volumes:
      - dbdata:/var/lib/mysql
    restart: always
    environment:
      MYSQL_ROOT_PASSWORD: somewordpress
      MYSQL_DATABASE: wordpress
      MYSQL_USER: wordpress
      MYSQL_PASSWORD: wordpress
  wordpress:
    depends_on:
      - db
    image: wordpress:latest
    ports:
      - "8000:80"
    restart: always
    environment:
      WORDPRESS_DB_HOST: db:3306
      WORDPRESS_DB_USER: wordpress
      WORDPRESS_DB_PASSWORD: wordpress
volumes:
  dbdata:
```

Docker Compose

- Desplegar/arrancar los contenedores

```
docker-compose up
```

- Parar los contenedores

```
docker-compose stop
```

- Eliminar contenedores + red

```
docker-compose down
```

- Eliminar contenedores + red + volúmenes

```
docker-compose down --volumes
```


Docker Compose

- “Construir” o “Reconstruir” las imágenes

```
docker-compose build
```

- Reiniciar los contenedores

```
docker-compose restart
```

- Descargar las imágenes necesarias

```
docker-compose pull
```

- Modo “detached” (no se muestran los logs)

```
docker-compose --detach
```

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:
 - Instalar Docker (en Almalinux):
 - `sudo yum install -y yum-utils device-mapper-persistent-data lvm2`
 - `sudo yum-config-manager --add-repo`
 - `sudo yum install docker-ce`
 - Inicializar Docker:
 - `sudo systemctl start docker`

Docker Compose: cluster de pruebas

■ Virtualizar un cluster con Docker compose:

□ Instalar Docker Compose (en Almalinux):

■ Descargar de Github:

```
sudo curl -L
"https://github.com/docker/compose/releases/download/1.27.4/docker-compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
```

■ Le assignamos permisos de ejecución:

```
sudo chmod +x /usr/local/bin/docker-compose
```

```
[root@localhost ~]# docker-compose --version
docker-compose version 1.27.4, build 40524192
[root@localhost ~]#
```

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:
 - Instalar el cluster Big Data:
 - Descargar las imágenes de Hadoop de Big Data Europe:
<https://github.com/big-data-europe/docker-hadoop>
 - Descomprimir las imagenes: `unzip docker-hadoop-master.zip`
 - Entrar en la carpeta: `cd docker-hadoop-master`
 - Desplegar el cluster de 3 nodos:
`docker-compose up -d`

```
Creating historyserver    ... done
Creating nodemanager     ... done
Creating datanode        ... done
Creating resourcemanager ... done
Creating namenode        ... done
[root@localhost docker-hadoop-master]#
```


Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:
 - Comprobar que se desplegaron los contenedores

```
[root@localhost docker-hadoop-master]# docker ps
```

CONTAINER ID	STATUS	IMAGE	PORTS	COMMAND	NAMES	CREATED
1399da4bfa7e	Up 9 minutes (healthy)	bde2020/hadoop-nodemanager:2.0.0-hadoop3.2.1-java8	8042/tcp	"/entrypoint.sh /run..."	nodemanager	9 minutes ago
e4ef316095a7	Up 9 minutes (healthy)	bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8	9864/tcp	"/entrypoint.sh /run..."	datanode	9 minutes ago
a8516291d3bf	Up 9 minutes (healthy)	bde2020/hadoop-resourcemanager:2.0.0-hadoop3.2.1-java8	8088/tcp	"/entrypoint.sh /run..."	resourcemanager	9 minutes ago
88c9bce49fea	Up 9 minutes (healthy)	bde2020/hadoop-historyserver:2.0.0-hadoop3.2.1-java8	8188/tcp	"/entrypoint.sh /run..."	historyserver	9 minutes ago
ddc27b9470d0	Up 9 minutes (healthy)	bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8	0.0.0.0:9000->9000/tcp, 0.0.0.0:9870->9870/tcp	"/entrypoint.sh /run..."	namenode	9 minutes ago

```
[root@localhost docker-hadoop-master]#
```



Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:
 - Ejecutar programas en el cluster:
 - Conectarse a un nodo: `sudo docker exec -it namenode bash`
 - Crear datos de entrada:

```
mkdir libros
cd libros
curl http://www.gutenberg.org/cache/epub/2000/pg2000.txt >
quijote.txt
curl http://www.gutenberg.org/cache/epub/345/pg345.txt >
dracula.txt
```
 - Cargar los datos en el Sistema de archivos distribuidos HDFS:

```
hdfs dfs -mkdir -p librosEnHDFS
hdfs dfs -put ./librosEnHDFS
```

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:

- Ejecutar programas en el cluster:

- Ejecutar programa:

```
hadoop jar
$HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-
examples-${HADOOP_VERSION}.jar wordcount
/user/root/librosEnHDFS/libros/
/user/root/salidaEnHDFS
```

Job Counters


```
Launched map tasks=2
Launched reduce tasks=1
Rack-local map tasks=2
Total time spent by all maps in occup
Total time spent by all reduces in oc
```

Map-Reduce Framework

```
Map input records=53834
Map output records=548684
Map output bytes=5204173
Map output materialized bytes=263653
Input split bytes=254
Combine input records=548684
Combine output records=59084
Reduce input groups=57977
Reduce shuffle bytes=263653
Reduce input records=59084
Reduce output records=57977
```

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:
 - Ejecutar programas en el cluster:
 - Comprobar la salida:
 - Desde servidor web:





Browse Directory

/user/root/salidaEnHDFS

Go!

Show 25 entries

Search:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	0 B	Nov 23 19:27	3	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	621.63 KB	Nov 23 19:27	3	128 MB	part-r-00000	

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:

- Ejecutar programas en el cluster:

- Comprobar la salida:

- Desde la CLI

```
root@ddc27b9470d0:/libros# hadoop fs -cat salidaEnHDFS/part-r-000000 | head
2020-11-23 18:33:46,175 INFO sasl.SaslDataTransferClient: SASL encryption trust
check: localhostTrusted = false, remoteHostTrusted = false
!Mal      1
"'Are     1
"'E's     1
"'I       1
"'Ittin'   1
"'Little   1
"'Lucy,    1
"'Maybe   1
"'Miss     1
"'My       2
cat: Unable to write to output stream.
root@ddc27b9470d0:/libros#
```

Docker Compose: cluster de pruebas

- Virtualizar un cluster con Docker compose:

- Apagar cluster:

- Salir del contenedor: Control + D

```
root@ddc27b9470d0:/libros# exit
[root@localhost docker-hadoop-master]#
```

- Apagar los contenedores: docker-compose down

```
[root@localhost docker-hadoop-master]# docker-compose down
Stopping nodemanager    ... done
Stopping datanode       ... done
Stopping resourcemanager ... done
Stopping historyserver  ... done
Stopping namenode       ... done
Removing nodemanager    ... done
Removing datanode       ... done
Removing resourcemanager ... done
Removing historyserver  ... done
Removing namenode       ... done
Removing network docker-hadoop-master_default
[root@localhost docker-hadoop-master]#
```

Monitorizar Docker

- Crear docker-compose con servicios:
 - ☐ cAdvisor: obtiene métrica del uso de contenedores
 - ☐ Prometheus: base de datos que almacena las métricas
 - ☐ caddy: servidor web y proxy reverso

Monitorizar Docker

docker-compose.yml

```
version: '3.9'
```

```
volumes:  
  prometheus_data: {}
```

```
services:  
  prometheus:  
    image: prom/prometheus:v2.33.3  
    container_name: prometheus  
    expose:  
      - 9090  
    command:  
      - --config.file=/etc/prometheus/prometheus.yml  
      - --storage.tsdb.retention.time=1y  
    volumes:  
      - ./prometheus.yml:/etc/prometheus/prometheus.yml:ro  
      - prometheus_data:/prometheus:rw  
    depends_on:  
      - cadvisor
```

```
  cadvisor:  
    image: gcr.io/cadvisor/cadvisor:v0.43.0  
    container_name: cadvisor  
    expose:  
      - 8080  
    volumes:  
      - /:/rootfs:ro  
      - /var/run:/var/run:rw  
      - /sys:/sys:ro  
      - /var/lib/docker:/var/lib/docker:ro  
    command:  
      - "--enable_load_reader=true"
```

```
  caddy:  
    image: caddy:2.4.6  
    container_name: caddy  
    ports:  
      - "9090:9090"  
      - "8080:8080"  
    volumes:  
      - ./Caddyfile:/etc/caddy/Caddyfile  
    environment:  
      - ADMIN_USER=${ADMIN_USER:-admin}  
      - ADMIN_PASSWORD=${ADMIN_PASSWORD:-admin}  
      - ADMIN_PASSWORD_HASH=${ADMIN_PASSWORD_HASH:-}
```

Monitorizar Docker

```
version: '3.9'
```

docker-compose.yml

```
volumes:  
  prometheus_data: {}
```

```
services:  
  prometheus:  
    image: prom/prometheus:v2.33.3  
    container_name: prometheus  
    expose:  
      - 9090
```

Para saber cuál es el hash de la contraseña que queremos asignar: escribimos la contraseña

```
jesus@injest:/disk0/monitorizar$ sudo docker run --rm caddy:2.4.6 caddy hash-password --plaintext ' 
```

```
volumes:  
  - ./prometheus.yml:/etc/prometheus/prometheus.yml:ro  
  - prometheus_data:/prometheus:rw  
depends_on:  
  - cadvisor
```

```
cadvisor:  
  image: gcr.io/cadvisor/cadvisor:v0.43.0  
  container_name: cadvisor  
  expose:  
    - 8080  
  volumes:  
    - /:/rootfs:ro  
    - /var/run:/var/run:rw  
    - /sys:/sys:ro  
    - /var/lib/docker:/var/lib/docker:ro  
  command:  
    - "--enable_load_reader=true"
```

```
caddy:  
  image: caddy:2.4.6  
  container_name: caddy  
  ports:  
    - "9090:9090"  
    - "8080:8080"  
  volumes:  
    - ./Caddyfile:/etc/caddy/Caddyfile  
  environment:  
    - ADMIN_USER=${ADMIN_USER:-admin}  
    - ADMIN_PASSWORD=${ADMIN_PASSWORD:-admin}  
    - ADMIN_PASSWORD_HASH=${ADMIN_PASSWORD_HASH:-
```

Copiamos el hash a la variable de entorno

Monitorizar Docker



Caddyfile



docker-compose.yml



prometheus.yml

■ Archivos:

- docker-compose.yml: contiene los contenedores
- prometheus.yml:
 - Cada 10s guarda las métricas de cadvisor

```
scrape_configs:  
  - job_name: cadvisor  
    scrape_interval: 10s  
    static_configs:  
      - targets:  
        - cadvisor:8080
```

Monitorizar Docker



Caddyfile



docker-compose.yml



prometheus.yml

■ Archivos:

- ☐ docker-compose.yml: contiene los contenedores
- ☐ prometheus.yml:
- ☐ Caddyfile: autenticación básica (usuario y contraseña) y proxy reverso a cadvisor y prometheus

```
:9090 {  
    basicauth /* {  
        {$ADMIN_USER} {$ADMIN_PASSWORD_HASH}  
    }  
    reverse_proxy prometheus:9090  
}  
  
:8080 {  
    basicauth /* {  
        {$ADMIN_USER} {$ADMIN_PASSWORD_HASH}  
    }  
    reverse_proxy cadvisor:8080  
}
```

Monitorizar Docker


■ Desplegar servicio de monitorización:

```
jesus@injtest:/disk0/monitorizar$ sudo docker-compose up --detach
Creating network "monitorizar_default" with the default driver
Creating volume "monitorizar_prometheus_data" with default driver
Creating cadvisor ... done
Creating caddy ... done
Creating prometheus ... done
jesus@injtest:/disk0/monitorizar$
```

- ☐ Crea red
- ☐ Crea volumen para guardar los datos de la bd
- ☐ Crea contenedores

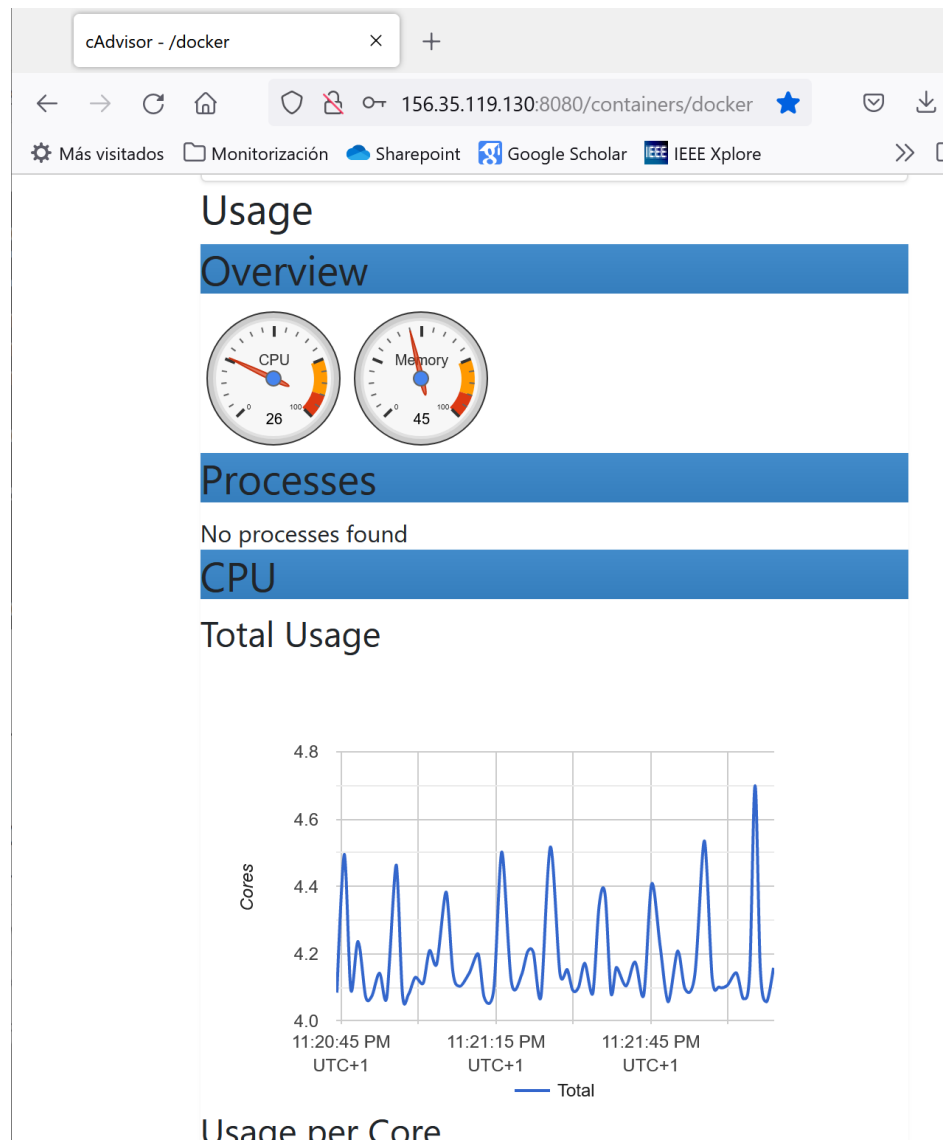
Monitorizar Docker

■ cAdvisor




Monitorizar Docker

■ cAdvisor



Monitorizar Docker

■ Prometheus:



Data + Dev + Op

■ Devops:



- ☐ Desarrollo + Operaciones
- ☐ Integración continua
- ☐ Despliegue continuo
- ☐ Entrega continua

■ DataOps

- ☐ Desarrollo + Operaciones + científicos/analistas de datos
- ☐ Analítica continua

Orquestación de contenedores

- Organizar y dirigir cómo se despliegan los contenedores. Alternativas:



"Kubernetes" by laboratoriolinux is licensed under CC BY-NC-SA 2.0



Jesús Morán y Cristian Augusto

Grupo de Investigación en Ingeniería del Software

<http://giis.uniovi.es>

Universidad de Oviedo

